

# Introduction to Database Systems: CS305

## Advanced queries, joins and aggregates

Oliver Bonham-Carter  
Hang Zhao

17-19 October 2023

# Joins: Bringing Data Together

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao



## Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

## Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- The SQLite3 join-clause is used to combine records from two or more tables in a database.
- A **JOIN** is a means for combining fields from two tables by using values common to each.

# Joins: Visual Definitions

## As Venn Diagrams

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

### Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

### Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

Inner Join




Table 1

Table 2

Left Join




Table 1

Table 2

Right Join




Table 1

Table 2

Full Outer Join




Table 1


Table 2

# SQL Code and Venn Diagrams


Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao


## SQL JOINS




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```




```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

### Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

### Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

# Joins: Visual Definitions

## Combining Tables

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates



Oliver  
Bonham-  
Carter  
Hang Zhao

### Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

### Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having



# An Explanation of Terms

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms

Inner Join

Left Join

Right Join

Cross Joins

Fine Tuning

Sets

AS Clauses

Strings

Ordering

Having

- SQL joins

- An **inner join** will return records that have matching values in *both* tables.
- A left **outer join** will return all records from the *left* table and the matched records from the *right* table.
- A right **outer join** will return all records from the *right* table and the matched records from the *left* table.
- A **full outer join** will return all records when there is a match from *either* table.

# Inner Joins

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins


Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

Student ID	Name	Student ID	Department
1001	A	1004	Mathematics
1002	B	1005	Mathematics
1003	C	1006	History
1004	D	1007	Physics
		1008	Computer Science

Table A

Table B



Student ID	Name	Department
1004	D	Mathematics

# Inner joins

File: /sandbox/fruitJoin.txt

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

## Create two tables

```
DROP TABLE IF EXISTS TableA;  
CREATE TABLE TableA (  
  fruit VARCHAR,  
  colour VARCHAR);
```

```
DROP TABLE IF EXISTS TableB;  
CREATE TABLE TableB (  
  fruit VARCHAR,  
  colour VARCHAR);
```



# Inner joins

File: /sandbox/fruitJoin.txt

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

## Populate the tables

```
INSERT INTO TableA VALUES ("Lemons_A","Yellow");  
INSERT INTO TableA VALUES ("Apples_A","Red");  
INSERT INTO TableA VALUES ("Grapes_A","Purple");
```

```
INSERT INTO TableB VALUES ("Lemons_B","Yellow");  
INSERT INTO TableB VALUES ("Apples_B","Red");  
INSERT INTO TableB VALUES ("Oranges_B", "Orange");  
INSERT INTO TableB VALUES ("Grapes_B","Purple");
```

# Inner joins

File: /sandbox/fruit - innerJoin.txt

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

## Use INNER JOIN to query

```
.tables
```

```
SELECT * from TableA;
```

```
SELECT* from TableB;
```

```
SELECT
```

```
    TableA.fruit,  
    TableA.colour,  
    TableB.colour,  
    TableB.fruit
```

```
FROM
```

```
    TableA
```

```
INNER JOIN
```

```
    TableB ON TableB.colour == TableA.colour;
```

# Inner joins

## Output

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

## Output

```
Lemons_A|Yellow|Yellow|Lemons_B
Apples_A|Red|Red|Apples_B
Grapes_A|Purple|Purple|Grapes_B
```

# Left Join

Matches entries from LEFT table to the other table

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
**Left Join**  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

## Setup Tables

```
/*Drop the table if it already exists*/  
DROP TABLE IF EXISTS Employees;  
  
/*Create the Employees table*/  
CREATE TABLE Employees (  
    EmployeeID INT PRIMARY KEY,    FirstName VARCHAR,  
    LastName VARCHAR,    DepartmentID INT  
);  
  
DROP TABLE IF EXISTS Departments;  
  
CREATE TABLE Departments (  
    DepartmentID INT PRIMARY KEY,    DepartmentName VARCHAR  
);
```

# Left Join

example

Introduction  
to Database  
Systems:  
CS305

Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

## Populate

```
/*Insert sample data into the Employees table*/
```

```
INSERT INTO Employees (  
    EmployeeID, FirstName,  
    LastName, DepartmentID)
```

```
VALUES
```

```
    (1, 'John', 'Doe', 1),  
    (2, 'Jane', 'Smith', 2),  
    (3, 'Bob', 'Johnson', 1),  
    (4, 'Alice', 'Williams', NULL);
```

```
/*Insert sample data into the Departments table*/
```

```
INSERT INTO Departments (DepartmentID, DepartmentName)
```

```
VALUES
```

```
    (1, 'HR'),  
    (2, 'IT'),  
    (3, 'Finance');
```

# Left Join

example

## Query

```
/*Perform a LEFT JOIN to retrieve a list
of all employees and their departments*/
SELECT
    e.EmployeeID, e.FirstName, e.LastName, d.DepartmentName
FROM
    Employees e
LEFT JOIN
    Departments d
ON
    e.DepartmentID = d.DepartmentID;
```

EmployeeID	FirstName	LastName	DepartmentName
1	John	Doe	HR
2	Jane	Smith	IT
3	Bob	Johnson	HR
4	Alice	Williams	NULL

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

# Right Join

Matches entries from RIGHT table to the other table

## Setup Tables

```
/*Perform a LEFT JOIN to retrieve a list
  of all employees and their departments*/
SELECT
    e.EmployeeID, e.FirstName, e.LastName, d.DepartmentName
FROM
    Employees e
RIGHT JOIN
    Departments d
ON
    e.DepartmentID = d.DepartmentID;
```

EmployeeID	FirstName	LastName	DepartmentName
1	John	Doe	HR
2	Jane	Smith	IT
3	Bob	Johnson	HR
NULL	NULL	NULL	Finance

# Cross joins

## Cross product derived from both tables

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

```
DROP TABLE IF EXISTS ranks;
CREATE TABLE ranks (
    rank TEXT NOT NULL
);
DROP TABLE IF EXISTS suits;
CREATE TABLE suits (
    suit TEXT NOT NULL
);

INSERT INTO ranks(rank)
VALUES('2'),('3'),('4'),('5'),('6'),('7'),('8'),('9'),('10'),('J'),('Q'),('K'),('A');

INSERT INTO suits(suit) VALUES('Clubs'),('Diamonds'),('Hearts'),('Spades');

SELECT rank, suit
FROM ranks
CROSS JOIN suits
ORDER BY suit;
```



# Cross joins: All Card Pairs

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

## Cross join: output

2 | Clubs

...

J | Clubs

Q | Clubs

K | Clubs

A | Clubs

...

A | Spades



# New Database

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having



(A New Database!)

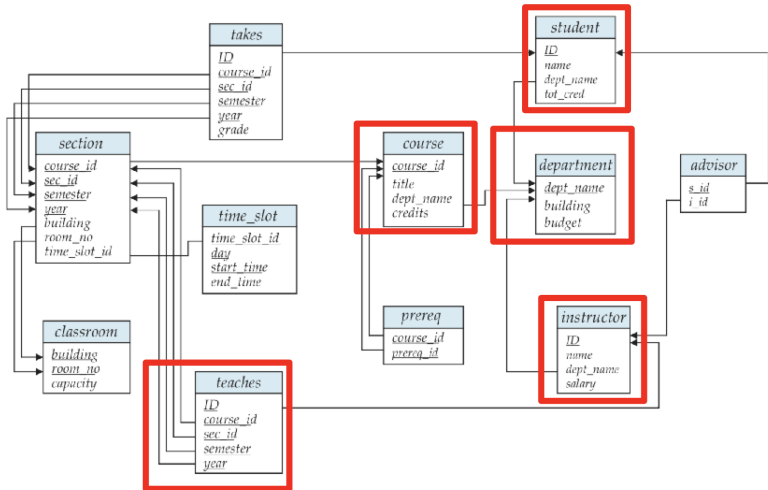


Schema: Red boxes are the tables of today's database study

Schema: Red boxes are the tables of today's database study

## Fine Tuning

- Sets
- AS Clauses
- Strings
- Ordering
- Having



**Figure 2.8** Schema diagram for the university database.

# New Database

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- Find the database maker file, *campusDB\_build.txt*, in your sandbox directory

```
cat campusDB_build.txt | sqlite3 myCampusDB.sqlite3
```

# Set Operations

## OR & AND

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- **OR:** Find all deptNames in the **UNION** of Instructor and Course

- `select deptName from Instructor UNION select deptName from course;`

- `select distinct(deptName) from Instructor;`

- **AND:** Find all deptNames in the **INTERSECT** of Instructor and Course

- `select deptName from Instructor INTERSECT select deptName from Course;`

- `select distinct(Instructor.deptName) from Instructor, Course where Instructor.deptName == Course.deptName;`

# Set Operations

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- `select distinct(deptName) from Instructor;`
- `select distinct(deptName) from Course;`

- The EXCEPT operator compares the result sets of two queries and returns distinct rows from the left query that are not in the output by the right query.
- Find all deptNames different to both the *Instructor* and *Course*
- Check these two queries below. Why is the output different?

- `select deptName from Instructor EXCEPT select deptName from Course;`
- `select deptName from Course EXCEPT select deptName from Instructor;`

- The **AS** clause is used to rename relations; useful for reducing necessary code in queries
- Ex: *For all instructors in the university who have taught some course, find their names and the course ID of all their taught courses*
  - Select **I.name**, **T.courseID**  
FROM **Instructor AS I** , **Teaches AS T**  
WHERE **I.ID= T.ID**;
- On the second line:
  - the **Instructor** table is renamed to **I**
  - the **Teaches** table is renamed to **T**.

- Another reason to rename a relation is a case where we wish to compare tuples in the same relation.
- We then need to take the Cartesian product of a relation with itself and, without renaming, it becomes impossible to distinguish one tuple from the other.
- Suppose that we want to write the query, *find the names of all instructors whose salary is greater than at least one instructor in the Math department.*
  - ```
SELECT DISTINCT T.name
FROM Instructor as T ,
Instructor AS S
WHERE T.salary > S.salary and S.deptName == "Math"
```



# AS clauses

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
**AS Clauses**  
Strings  
Ordering  
Having

- Find all names of common teachers in Instructor and Teaches tables

## Use AS to implement variables attributes to hold places

- `select distinct(Instructor.name) as newName from Instructor, teaches where Instructor.ID = teaches.ID and newName == "Thompson";`

- Find the names of all Instructors whose salary is greater than at least one Instructor in the Math department.
- `select distinct(T.name) from Instructor as T,  
Instructor as S where T.salary > S.salary and  
S.deptName == "Math";`
- `select distinct T.name, T.salary from Instructor as  
T, Instructor as S where T.salary > S.salary and  
S.deptName == "Math";`
- Reference: `select * from Instructor;`

# Regular Expression-ish

- Textual *wildcards* to recover information from partial knowledge.
- Finding substrings using the % and \_ operators.

- `select name from Instructor where name like "%ille%";`
  - Selects *Miller* from a substring
- `select name from Instructor where name like "%son";`
  - Selects all names followed by "son" substring
- Compare to: *Select \* from Instructor;*
- `select name from Instructor where name like "__11__";`
- `select name from Instructor where name like "__11___";`
  - Selects "Miller" or "William" from the number of spaces after the "ll";

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

# Regular Expression-ish

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
**Strings**  
Ordering  
Having

- Find special pattern characters (i.e., "%" and "\_") in strings
- SQL even allows the specification of an escape character.

- like 'ab\%cd%' escape '\'** matches all strings beginning with "ab%cd".
- like 'ab\\cd%' escape '\'** matches all strings beginning with "ab\cd".

# Ordering Results

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
**Ordering**  
Having

- SQL allows for sorting the output.
- Output is sorted alphabetically

- `select name from Instructor order by name;`
- `select name,salary from Instructor order by salary;`
  - Provides numerical values in an interval

# “Intermediate” Results Using **HAVING**

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

- The **HAVING** clause enables you to specify conditions that filter which group results appear in the final results.
- The **HAVING** clause must follow the GROUP BY clause in a query and must also precede the ORDER BY clause if used.

## Pseudo-code

```
SELECT column1, column2
FROM table1, table2
WHERE [ conditions ]
GROUP BY column1, column2
HAVING [ conditions ]
ORDER BY column1, column2
```

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

# Group By

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- Give the number of names, and names of all members of departments who make less than 100000.

- `select count(name), deptName from Instructor GROUP BY deptName HAVING salary < 100000;`

- Give the deptNames and the average salaries for departments that begin with the letter 'C'.

- `select deptName, avg(salary) from Instructor group by deptName HAVING deptName LIKE "C%";`

# Group By

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- Give the department names and salaries from the Instructor group for whose members make between 97K and 100K.

- `select deptName, salary from Instructor group by deptName HAVING salary < 100000 and salary > 97000;`

- Compare to: Give me deptName and salary information where the salary is between 97K and 100K.

- `select deptName, salary from Instructor where salary < 100000 and salary > 97000 group by deptName;`



# Use avg to Query

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- `select deptName, avg(salary) from Instructor group by deptName;`
  - Report average salaries for departments
- `select deptName, avgSalary FROM (select deptName, avg(salary) as avgSalary from Instructor group by deptName) where avgSalary > 97000;`
  - Report average salaries larger than \$97k. This query is similar to one using the **HAVING** clause. Here we use the **FROM** clause.

# Ordering Result Using BETWEEN

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao

Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- SQL allows for sorting the output by criteria
- Output is sorted for values in an interval

- `select name, salary from Instructor where salary <= 100000 and salary >= 90000;`
- `select name, salary from Instructor where salary between 70000 and 100000;`
  - Query values in their intervals.

# Consider this!

Introduction  
to Database  
Systems:  
CS305  
Advanced  
queries, joins  
and  
aggregates

Oliver  
Bonham-  
Carter  
Hang Zhao



Joins

Terms  
Inner Join  
Left Join  
Right Join  
Cross Joins

Fine Tuning

Sets  
AS Clauses  
Strings  
Ordering  
Having

- Can you create a JOIN between two tables?
- Can you use EXCEPT and AS to add fine tune your queries?
- Can you write SQL code to be more precise numerically using BETWEEN, AVG, and greater-than and less-than?